




# Container & K8S


- 컴퓨터(가상머신)와 컨테이너 차이




Container





Container




Container




Container



Container



Container




Container




Container


# Single Host Docker Platform (WebServer 구축, PaaS)




Docker Hub(hub.test.com)



#docker push hub.test.com/tomcat  
#docker push hub.test.com/apache  
#docker push hub.test.com/mysql



# Multi Host Docker Platform (쇼핑몰 구축 PaaS)




Docker Hub(hub.test.com)




① nginx image push


#docker push hub.test.com/nodejs  
#docker push hub.test.com/python  
#docker push hub.test.com/tomcat



# Multi Host Docker Platform (HA 구성)



Docker Hub(hub.test.com)




① nginx image push


#docker push hub.test.com/nodejs

#docker push hub.test.com/python


#docker push hub.test.com/tomcat




# Multi Host Docker Platform (쇼핑몰 구축 PaaS)



Docker Hub(hub.test.com)



#docker push hub.test.com/basket.js  
#docker push hub.test.com/pay.py  
#docker push hub.test.com/tomcat  
#docker push hub.test.com/mysql



# Docker & K8S

- 컨테이너 오케스트레이션 : 컨테이너의 배포, 관리, 확장, 네트워킹을 자동화

Layer 5	Orchestration Scheduling	<b>Kubernetes(K8S)</b> , Docker Swarm
Layer 3~4	Container Engine	<b>Docker</b> , Rocket, LxC, LxD
	Operating System	<b>Ubuntu</b> , CentOS, Fedora, RHEL
Hardware Layer (Layer 1~2)	Virtual Infrastructure	Openstack, vSphere, Azure
	Physical Infrastructure	Computer, Network, Storage

# Kubernetes(k8s)


- 배의 조타수란 그리스어에서 유래
- 2014년 구글에서 개발
- Kubernetes== k8s
- **컨테이너 orchestration(오케스트레이션)**
  - 관현악 편곡, 통합, 결집, 편성, 조직화, 조정
  - 운영 수준의 컨테이너 오케스트레이션
  - 컨테이너와 호스트의 수가 증가함에 따라 여러 컨테이너의 배포 프로세스를 최적화 시킴
  - 컨테이너화 된 애플리케이션을 자동으로 배포, 스케일링 및 관리해 주는 오픈 소스 시스템



# K8s 기능

- 컨테이너가 수천 개, 수만 개, 수십억 개가 된다면 관리가 복잡하고, 어려워짐
- 컨테이너를 통합·관리해야 하는 필요에 의해 등장한 기술
- 쿠버네티스를 사용하는 조직은 리눅스 기반의 컨테이너화된 애플리케이션들을 실행하는 호스트 그룹을 묶어 클러스터링을 구성해 사용하는 방식을 채용하고 있음
- 쿠버네티스는 이렇게 구성된 클러스터를 관리
- 또한 클러스터 내 애플리케이션 컨테이너들을 자동으로 배치하고 스케일링하는 등 여러 운영 작업을 자동화함
- 구글 ~ 2020년 기준 G메일, 구글 드라이브 등 애플리케이션을 구동하기 위해 쿠버네티스를 활용, 쿠버네티스로 약 30억 개의 컨테이너를 운영 중
- 국내 대표 스타트업 10곳 중 7곳은 이미 수백~수천 개의 컨테이너를 운영 중
- 깃허브, 카카오톡, 다음, 멜로, VISA 등 다수의 기업이 K8S 기반 환경을 구축해서 사용 중

## (예) 2대의 웹서버가 필요한 경우



# 명령어 kubectl

# kubectl

- k8s에게 원하는 작업을 요청 시 사용하는 명령어
- k8s cluster를 관리하는 동작은 kubectl이라는 Command line interface로 실행
  - K8s 자원들의 생성, 업데이트, 삭제 (create, update, delete)
  - 디버그, 모니터링, 장애처리(log, exec, cp, top, attach..)
  - 클러스터 관리(cordon, top, drain, taint...)

kubectl --help

```
root@masternode:~# kubectl --help
kubectl controls the Kubernetes cluster manager.

Find more information at: https://kubernetes.io/docs/reference/kubectl/ove

Basic Commands (Beginner):
  create      Create a resource from a file or from stdin
  expose      Take a replication controller, service, deployment or pod a
  run         Run a particular image on the cluster
  set         Set specific features on objects

Basic Commands (Intermediate):
  explain     Get documentation for a resource
  get         Display one or many resources
  edit        Edit a resource on the server
  delete      Delete resources by file names, stdin, resources and names,
```

## kubectl 명령어 형식

kubectl [command] [TYPE] [NAME] [flags]

<b>Command</b>	자원에 실행되는 동작	create, get, delete
<b>TYPE</b>	자원타입	pod, service, ingress
<b>NAME</b>	자원이름	
<b>Flags</b>	부가적으로 설정할 옵션	--help, -o wide

(ex) kubectl get pod WEBServer -o wide

➔ WEBServer 이름을 가진 Pod 자원정보를 자세히 확인

```
kubectl --help
```

```
kubuctl run --help
```

```
root@masternode:~# kubectl run --help
Create and run a particular image in a pod.

Examples:
# Start a nginx pod
kubectl run nginx --image=nginx

# Start a hazelcast pod and let the container expose port 5701
kubectl run hazelcast --image=hazelcast/hazelcast --port=5701

# Start a hazelcast pod and set environment variables "DNS_DOMAIN=clus
container
kubectl run hazelcast --image=hazelcast/hazelcast --env="DNS_DOMAIN=cl


# Start a hazelcast pod and set labels "app=hazelcast" and "env=prod"
kubectl run hazelcast --image=hazelcast/hazelcast --labels="app=hazelc

# Dry run; print the corresponding API objects without creating them
kubectl run nginx --image=nginx --dry-run=client
```


# K8S Pod

# Pod(파드)

- 컨테이너들을 모은 애플리케이션의 최소 단위
- K8s는 Pod 단위로 컨테이너들을 묶어서 관리
- 파드에는 하나 또는 여러 개의 컨테이너가 포함 될 수 있음



- K8s에는 container를 생성하는 API가 없음
- Pod API를 통해 pod를 생성/동작 시키고 pod를 통해 container를 생성/동작 시킴



## K8S에서 Pod 생성 방법 2가지

- 명령어 kubectl run으로 pod 생성과 실행
- YAML(야물) 파일을 이용한 pod 생성과 실행

**kubectl create ~ 또는 kubectl run ~**

## 1) kubectl run으로 pod 생성


- kubectl run web --image=nginx:1.14 --port 80
- kubectl get pods
- kubectl get pods -o wide

```
root@master:/k8s# kubectl run web --image=nginx:1.14 --port 80
pod/web created
root@master:/k8s#
root@master:/k8s# kubectl get pod -o wide
NAME    READY    STATUS    RESTARTS   AGE      IP           NODE    NOMINATED NODE   READINESS GATES
web     1/1     Running   0          45s     192.168.19.70   worker03 <none>        <none>
```

## 2) YAML 파일을 이용한 pod 생성과 실행

- YAML(야믈, *YAML ain't markup language*)
    - 마크업 언어가 아니다.
    - 직렬언어임
  - YAML 문법
    - 들여쓰기에 따라 구조가 바뀜
    - 들여쓰기는 **Tab이 아닌 Space bar 사용** (공백 문자들로 기본구조를 구성)
    - #로 주석을 표시
- ① Scalars(string/numbers) : ‘Key: value’ 형태로 표시  
‘:’ 다음에 공백 문자가 와야한다. (붙여쓰면 오류)
- ② Sequence(arrays/lists) : ‘-’(하이픈) 여러 개 나열

## 실습1. Single Pod



Pod : nginx-pod

```
#watch kubectl get pods -o wide  
#kubectl create -f pod-nginx.yaml
```

```
root@master:/K8s# cat pod-nginx.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx-pod  
spec:  
  containers:  
    - name: nginx-container  
      image: nginx:1.14  
      ports:  
        - containerPort: 80  
          protocol: TCP  
root@master:/K8s#
```

```
#kubectl describe pod pod-nginx
```

```
apiVersion: v1          //사용하려는 K8S API 버전 명시
kind: Pod               //어떤 종류의 object 또는 controller에 작업인지 명시
metadata:              //메타데이터 설정(Pod명)
  name: nginx-pod
Spec:                  //파드가 어떤 컨테이너를 갖고 실행하며 실행 시 어떻게 동작해야 할지 명시
  containers:
    - name: nginx-container
      image: nginx:1.14
      ports:
        - containerPort: 80
          protocol:tcp
```


*Nginx-pod*



*nginx-contaier*

```
#kubectl get pod -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
nginx-pod	1/1	Running	0	12m	192.168.19.67	worker03	<none>	<none>



## Pod안 Container 운영체제 확인




```
#kubectl exec nginx-pod -c nginx-container -it -- /bin/bash  
cat /etc/os-release
```

```
root@master:/k8s# kubectl exec nginx-pod -c nginx-container -it -- /bin/bash  
root@nginx-pod:# cat /etc/os-release  
PRETTY_NAME="Debian GNU/Linux 9 (stretch)"  
NAME="Debian GNU/Linux"  
VERSION_ID="9"  
VERSION="9 (stretch)"  
ID=debian  
HOME_URL="https://www.debian.org/"  
SUPPORT_URL="https://www.debian.org/support"  
BUG_REPORT_URL="https://bugs.debian.org/"  
root@nginx-pod:#
```

## Pod안 Container 접근 후 내용 수정

```
#kubectl exec nginx-pod -c nginx-container -it -- /bin/bash  
cd /usr/share/nginx/html  
echo "Hello~ Web" > index.html  
cat index.html  
exit  
#curl 192.168.19.67
```

```
root@master:/k8s# kubectl exec nginx-pod -c nginx-container -it -- /bin/bash  
root@nginx-pod:# cd /usr/share/nginx/html  
root@nginx-pod:/usr/share/nginx/html# ls  
50x.html  index.html  
root@nginx-pod:/usr/share/nginx/html# echo "Hello~Web">>index.html  
root@nginx-pod:/usr/share/nginx/html# cat index.html  
Hello~Web  
root@nginx-pod:/usr/share/nginx/html# exit  
exit  
root@master:/k8s# curl 192.168.19.67  
Hello~Web  
root@master:/k8s#
```



## 실습 명령어 정리

### ① Multi pod 생성

nano pod-nginx.yaml

kubectl create -f pod-nginx.yaml

### ② Pod 정보 확인

kubectl get pods

kubectl get pods -o wide

### ③ Pod 상세 정보 확인(장애처리)와 수정


kubectl describe pod pod-nginx

### ④ Pod 삭제

kubectl delete pod pod-nginx

kubectl delete pod --all

## 실습2. Multi-pod



Pod : multipod


```
#kubectl create -f pod-multi.yaml  
#kubectl get pods -o wide
```

```
root@master:/k8s# cat pod-multi.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
  name: multipod  
spec:  
  containers:  
    - name: nginx-container  
      image: nginx:1.14  
      ports:  
        - containerPort: 80  
    - name: ubuntu-container  
      image: ubuntu:24.04  
      command:  
        - sleep  
        - "10000"
```

```
#kubectl describe pod multipod
```

**watch kubectl get pod -o wide**


NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
multipod	0/2	ContainerCreating	0	8s	<none>	worker03	<none>	<none>	



**kubectl get pod -o wide**

root@master:/k8s# kubectl get pod -o wide									
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES
multipod	2/2	Running	0	12m	192.168.19.68	worker03	<none>	<none>	

## kubectl describe pod multipod




```
Node:           worker03/192.168.10.40
Start Time:    Tue, 13 Jan 2026 12:53:49 +0900
Labels:         <none>
Annotations:   cni.projectcalico.org/containerID: 9da97d6cdf1273d7e8
               cni.projectcalico.org/podIP: 192.168.19.68/32
               cni.projectcalico.org/podIPs: 192.168.19.68/32
Status:        Running
IP:            192.168.19.68
IPs:           IP: 192.168.19.68
Containers:
  nginx-container:
    Container ID: containerd://401fb51084fa56f9b0575d798f5e9b077f7cf2
    Image:         nginx:1.14
    Image ID:     docker.io/library/nginx@sha256:f7988fb6c02e0ce69257
    Port:          80/TCP
    Host Port:    0/TCP
    State:        Running
    Started:     Tue, 13 Jan 2026 12:53:50 +0900
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-acces
  ubuntu-container:
    Container ID: containerd://f7c6c35d8dbb421f7cc801131b1b30797d8e74a
    Image:         ubuntu:24.04
    Image ID:     docker.io/library/ubuntu@sha256:c35e29c9450151419d94
```

## kubectl describe pod multipod

Events:				
Type	Reason	Age	From	Message
Normal	Scheduled	5m13s	default-scheduler	Successfully assigned default/multipod to worker03
Normal	Pulled	5m12s	kubelet	Container image "nginx:1.14" already present on machine
Normal	Created	5m12s	kubelet	Created container: nginx-container
Normal	Started	5m12s	kubelet	Started container nginx-container
Normal	Pulling	5m12s	kubelet	Pulling image "ubuntu:24.04"
Normal	Pulled	5m	kubelet	Successfully pulled image "ubuntu:24.04" in 11.657s (11.658s)
Normal	Created	5m	kubelet	Created container: ubuntu-container
Normal	Started	5m	kubelet	Started container ubuntu-container

## Pod안 Container 운영체제 확인




```
#kubectl exec multipod -c ubuntu-container -it -- /bin/bash  
#cat /etc/os-release
```


```
root@master:/k8s# kubectl exec multipod -c ubuntu-container -it -- /bin/bash  
root@multipod:/# cat /etc/os-release  
PRETTY_NAME="Ubuntu 24.04.3 LTS"  
NAME="Ubuntu"  
VERSION_ID="24.04"  
VERSION="24.04.3 LTS (Noble Numbat)"  
VERSION_CODENAME=noble  
ID=ubuntu  
ID_LIKE=debian  
HOME_URL="https://www.ubuntu.com/"  
SUPPORT_URL="https://help.ubuntu.com/"  
BUG_REPORT_URL="https://bugs.launchpad.net/ubuntu/"  
PRIVACY_POLICY_URL="https://www.ubuntu.com/legal/terms-and-policies/privacy-policy"  
UBUNTU_CODENAME=noble  
LOGO=ubuntu-logo  
root@multipod:/#
```

## Pod 안 Container 접근 후 내용 수정

```
#kubectl exec multipod -c nginx-container -it -- /bin/bash  
#cd /usr/share/nginx/html  
#echo "Hello~ Web" > index.html  
#cat index.html  
#exit  
#curl 192.168.19.68
```



```
root@master:/k8s# kubectl exec multipod -c nginx-container -it -- /bin/bash  
root@multipod:/# cd /usr/share/nginx/html  
root@multipod:/usr/share/nginx/html# ls  
50x.html index.html  
root@multipod:/usr/share/nginx/html# echo "My WebServer^^">index.html  
root@multipod:/usr/share/nginx/html# cat index.html  
My WebServer^^  
root@multipod:/usr/share/nginx/html# exit  
exit  
root@master:/k8s# curl 192.168.19.68  
My WebServer^^  
root@master:/k8s#
```



## kubectl delete pods multiple 또는 kubectl delete pods --all

```
Every 2.0s: kubectl get pod -o wide

NAME      READY   STATUS        RESTARTS   AGE     IP           NODE    NOMINATED NODE   READINESS GATES
multipod  2/2     Terminating   0          27m    192.168.19.68  worker03 <none>        <none>

● 1 master × +  
root@master:/k8s# kubectl delete pods multipod
pod "multipod" deleted
```

```
Every 2.0s: kubectl get pod -o wide

No resources found in default namespace.

● 1 master × +  
root@master:/k8s# kubectl delete pods multipod
pod "multipod" deleted
root@master:/k8s#
```

## 실습 명령어 정리

### ① Multi pod 생성

vi pod-multi.yaml

kubectl create -f pod-multi.yaml

### ② Pod 정보 확인

kubectl get pods

kubectl get pods -o wide

### ③ Pod 상세 정보 확인(장애처리)와 수정

kubectl describe pod multipod

kubectl edit pod multipod

### ④ Pod 삭제

kubectl delete pod multipod

kubectl delete pod --all